

Message Passing

Shared memory vs Message Passing

- Shared memory systems:
 - there is a (common) shared address space throughout the system.
 - Communication takes place via shared data variables and control variables
 - Semaphores and monitors
- All multicomputer systems that do not have a shared address space.
- Hence, message Passing is used for communication.
- Shared memory systems are easy to develop and use
- Communication via message-passing can be simulated by communication via shared memory and vice-versa

Emulating message passing on a shared memory system

- The shared address space can be partitioned into disjoint parts, one part being assigned to each processor.
- “Send” and “receive” operations can be implemented by writing to and reading from the destination/sender processor’s address space, respectively.
- A separate location can be reserved as the mailbox for each ordered pair of processes.
- A P_i – P_j message-passing can be emulated by a write by P_i to the mailbox and then a read by P_j from the mailbox.
- Mailboxes can be assumed to have unbounded size.
- Synchronization primitives inform the receiver/sender after the data has been sent/received.

Emulating message passing on a shared memory system

- Each shared location can be modeled as a separate process;
 - “write” is emulated by sending update message to owner process;
 - “read” is emulated by sending query message to owner process
- Expensive operation
- Read and write operations are implemented by using network-wide communication
- Application can of course use a combination of shared memory and message-passing.
- Combination of shared memory system and message passing system can be used.
 - In multicomputer system, each “processor” may be a tightly coupled multiprocessor system with shared memory.
 - Between two computers, the communication is by message passing.

Primitives for distributed communication

- Broadly defined into two;
 - Blocking/non-blocking
 - Synchronous/asynchronous
- Message send and message receive communication primitives are denoted as:
 - **Send()** : at least two parameters – the destination, and the buffer in the user space, containing the data to be sent.
 - **Receive()**: at least two parameters – the source from which the data is to be received and the user buffer into which the data is to be received

Primitives for distributed communication

- There are two ways of sending data when the **Send primitive** is invoked
 - The buffered option : Copies data from the user buffer to the kernel buffer. The data later gets copied from the kernel buffer onto the network.
 - The unbuffered option: the data gets copied directly from the user buffer onto the network.
- For the **Receive primitive**, the buffered option is required
 - The data may already have arrived when the primitive is invoked
 - Needs a storage place in the kernel.

Primitives for distributed communication

- Synchronous primitives:
 - Both the Send() and Receive() handshake with each other.
 - The processing for the Send primitive completes only after the invoking processor learns that the corresponding Receive primitive has also been invoked and that the receive operation has been completed.
 - The processing for the Receive primitive completes when the data to be received is copied into the receiver's user buffer.
- Asynchronous primitives :
 - A Send primitive is asynchronous if control returns back to the invoking process after the data item to be sent has been copied out of the user-specified buffer.
 - It does not make sense to define asynchronous Receive primitives


Primitives for distributed communication

- Blocking primitives:
 - A primitive is blocking if control returns to the invoking process after the processing for the primitive (whether in synchronous or asynchronous mode) completes.
- Non-blocking primitives:
 - A primitive is non-blocking if **control returns back to the invoking process immediately after invocation**, even though the operation has not completed.
 - For a non-blocking Send, control returns to the process even before the data **is copied out of the user buffer**.
 - For a non-blocking Receive, control returns to the process even **before the data may have arrived from the sender**
 - For non-blocking primitives, **a return parameter** on the primitive call returns **a system-generated handle** which can be later used to check the status of completion of the call.
 - **Wait call.**


Primitives for distributed communication

- Blocking synchronous send
- Non-blocking synchronous send
- Blocking asynchronous send
- Non-blocking asynchronous send
- Blocking receive
- Non-blocking Receive


Primitives for distributed communication




(a) Blocking sync. *Send*, blocking *Receive*



(b) Nonblocking sync. *Send*, nonblocking *Receive*



(c) Blocking async. *Send*



(d) Non-blocking async. *Send*

- Duration to copy data from or to user buffer
- Duration in which the process issuing send or receive primitive is blocked

Primitives for distributed communication

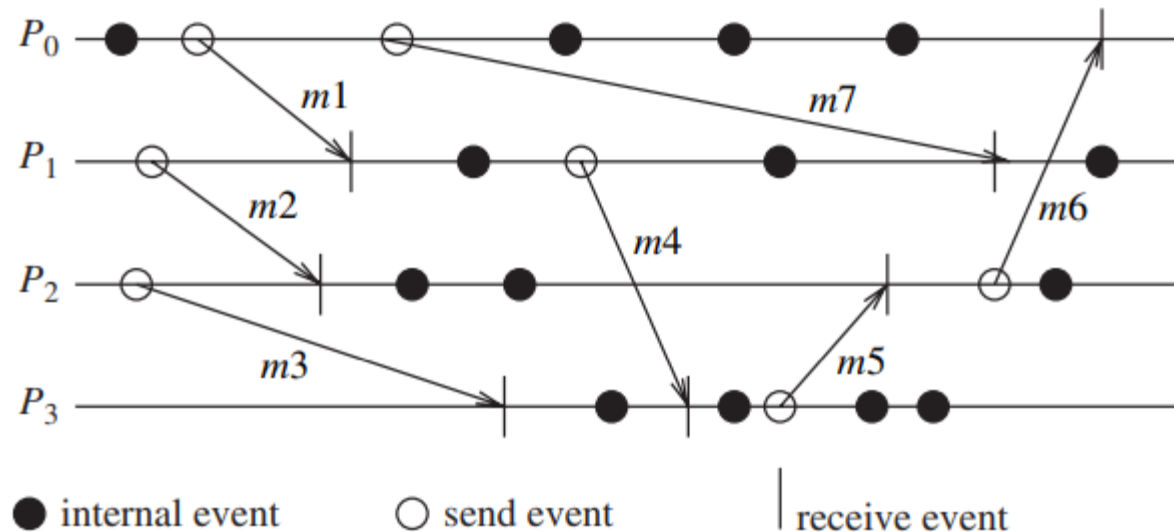
- A synchronous Send is **easier to use**.
- Handshake between the Send and the Receive makes the communication appear instantaneous, thereby **simplifying the program logic**.
- Synchronous Send **lowers the efficiency** within the sender process
- The non-blocking asynchronous Send is useful **for transferring large data**.
- Programmer has to **keep track of the completion** of primitive operations

Processor Synchrony

- Classification of **synchronous versus asynchronous processors**.
- Indicates that all the **processors execute in lock-step** with their clocks synchronized.
- As this synchrony is not attainable in a distributed system, a **large granularity of code** is termed as a step.
- For steps the processors are **synchronized**.
- Ensures that **no processor begins** executing the **next step** of code until all the processors **have completed executing the previous steps** of code.

Synchronous vs asynchronous executions


- An asynchronous execution is an execution in which :
 - (i) there is **no processor synchrony** and there is **no bound on the drift** rate of processor clocks,
 - (ii) message delays (transmission + propagation times) **are finite but unbounded**,
 - (iii) there is **no upper bound on the time** taken by a process to execute a step.



P_1 P_2
Com

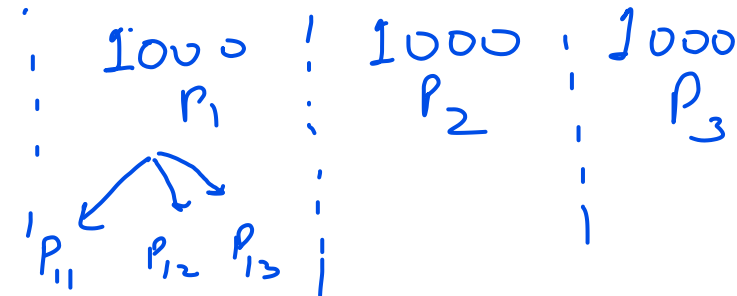
Synchronous vs asynchronous executions

- A synchronous execution is an execution in which :
 - processors are synchronized and the clock drift rate between any two processors is bound
 - message delivery (transmission + delivery) times are such that they occur in one logical step or round
 - there is a known upper bound on the time taken by a process to execute a step.



Synchronous vs asynchronous executions

- Easier to design and verify algorithms with synchronous executions.
- Practically difficult to build a completely synchronous system.
- Will involve delaying or blocking some processes for some time durations.



- **Virtually synchronous executions:**
 - Processors are allowed to have an asynchronous execution for a period of time and then they synchronize.
 - Within each round/phase/step, there may be a **finite and bounded** number of **sequential sub-rounds** (or subphases or sub-steps) that processor executes.
 - Each sub-round is assumed to **send at most one message per process**; hence the message(s) sent will reach in a single message hop.

Emulations

- An asynchronous program can be emulated on a synchronous system
 - Synchronous system is a special case of an asynchronous system
 - All communication finishes within the same round in which it is initiated
- A synchronous program can be emulated on an asynchronous system using a tool called synchronizer.

